

# Ngoc Bui

New Haven, CT, US · (+1) 203-809-4258

ngocbh.pt@gmail.com · ngoc.bui@yale.edu · ngocbh.github.io · scholar: LVk1D\_cAAAAAJ

## EDUCATION

---

- **Ph.D. in Computer Science** 2023 – 2028 [expected]  
*Yale University*
  - Advisor: Prof. Rex Ying
- **M.S. in Data Science** 2021 – 2023  
*Hanoi University of Science and Technology (HUST)*
  - GPA: 3.9/4.0, Major GPA: 4.0/4.0.
  - Thesis: Evolutionary and Deep Reinforcement Learning Algorithms for Optimizing the Lifetime of Wireless Sensor Networks.
  - Advisors: Prof. Thuan Do Phan and Dr. Phi Le Nguyen
- **Engineer in Computer Science** 2016 – 2021  
*Hanoi University of Science and Technology (HUST)*
  - GPA: 3.67/4.0, Major GPA: 3.88/4.0.
  - Honors: Excellent (≈ 1st class honors).

## UNDER REVIEW (MOST RECENT)

---

- Ngoc Bui, Hieu Trung Nguyen, Arman Cohan, Rex Ying. "Token Retention Is All You Need for Dynamic KV Cache Budgeting and Eviction in Vision-Language Models". *Under Review*
- Hiren Madhu, Ngoc Bui, Ali Maatouk, Leandros Tassiulas, Smita Krishnaswamy, Menglin Yang, Sukanta Ganguly, Kiran Srinivasan, Rex Ying. "HypRAG: Hyperbolic Dense Retrieval for Retrieval Augmented Generation". *Under Review*.
- Runjin Chen, Mingxuan Ju, Ngoc Bui, Dimosthenis Antypas, Stanley Cai, Xiaopeng Wu, Leonardo Neves, Neil Shah, Tong Zhao. "Enhancing Item Tokenization for Generative Recommendation through Self-Improvement". *Under Review*.

## PUBLICATIONS

---

- Ngoc Bui, Shubham Sharma, Simran Lamba, Saumitra Mishra, Rex Ying. "Cache What Lasts: Token Retention for Memory-Bounded KV Cache in LLMs". The *International Conference on Learning Representations (ICLR)*, 2026.
- Neil He, Jiahong Liu, Buze Zhang, Ngoc Bui, Ali Maatouk, Menglin Yang, Irwin King, Melanie Weber, Rex Ying. "Position: Beyond Euclidean - Foundation Models Should Embrace Non-Euclidean Geometries". **Oral Presentation** at *Learning on Graphs Conference (LOG)*, 2025. [arxiv].
- Ngoc Bui, Menglin Yang, Runjin Chen, Leonardo Neves, Mingxuan Ju, Rex Ying, Neil Shah, Tong Zhao. "Learning Along the Arrow of Time: Hyperbolic Geometry for Backward-Compatible Representation Learning". The *International Conference on Machine Learning (ICML)*, 2025. [arxiv].
- Ngoc Bui, Hieu Trung Nguyen, Shantanu Kumar, Julian Theodore, Weikang Qiu, Viet Anh Nguyen, Rex Ying. "Mixture-of-Personas Language Models for Population Simulation". The *Findings of the Association for Computational Linguistics (ACL Finding)*, 2025. [arxiv].
- Ngoc Bui, Hieu Trung Nguyen, Viet Anh Nguyen, and Rex Ying. "Explaining Graph Neural Networks via Structure-aware Interaction Index". The *International Conference on Machine Learning (ICML)*, 2024. [arxiv].
- Jiasheng Zhang, Jialin Chen, Ali Maatouk, Ngoc Bui, Qianqian Xie, Leandros Tassiulas, Jie Shao, Hua Xu, Rex Ying. "LitFM: A Retrieval Augmented Structure-aware Foundation Model For Citation Graphs". *ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD)*, 2025. [arxiv].

- **Ngoc Bui**, Duy Nguyen, Man-Chung Yue, and Viet-Anh Nguyen. “Coverage-Validity-Aware Algorithmic Recourse”. *The Operations Research (OPRE) Journal*, 2024. [arxiv].
- Duy Nguyen, **Ngoc Bui**, Viet-Anh Nguyen. “Feasible Recourse Plan via Diverse Interpolation”. *The International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2023. [arxiv].
- Duy Nguyen, **Ngoc Bui**, Viet-Anh Nguyen. “Distributionally Robust Recourse Action”. *The International Conference on Learning Representations (ICLR)*, 2023. **Honorable Mention in INFORMS Undergraduate Operations Research Prize**. [arxiv].
- **Ngoc Bui**, Duy Nguyen, Viet-Anh Nguyen. “Counterfactual Plans under Distributional Ambiguity”. *The International Conference on Learning Representations (ICLR)*, 2022. [arxiv].
- Tuan-Duy Hien Nguyen, **Ngoc Bui**, Duy Nguyen, Man-Chung Yue, Viet Anh Nguyen. “Robust Bayesian Recourse”. *The Association for Uncertainty in Artificial Intelligence (UAI)*, 2022. [arxiv].
- **Ngoc Bui**, Phi Le Nguyen, Viet Anh Nguyen, Phan Thuan Do. “A Deep Reinforcement Learning-based Adaptive Charging Policy for WRSNs”. *IEEE International Conference on Mobile Ad-Hoc and Smart Systems (MASS)*, 2022.
- **Ngoc Bui**, Tam Nguyen, Binh Huynh Thi Thanh, and Trong Vinh Le. “A phenotype-based multi-objective evolutionary algorithm for maximizing lifetime in wireless sensor networks with bounded hop”. *The Journal of Soft Computing*, 2023.
- **Ngoc Bui** and Viet-Trung Tran. “A Novel Conditional Random Fields Aided Fuzzy Matching in Vietnamese Address Standardization”. *The International Symposium on Information and Communication Technology (SoICT)*, 2019.

## EXPERIENCES

---

- **JPMorgan Chase & Co.** June 2025 - August 2025  
*Data & Analytics Research Associate*  
 - Developing an efficient inference algorithm to improve the memory and computation bottleneck of long-context and long-generation in LLM inference, providing  $\times 2$  inference speed and preserving (sometimes outperforming) the performance with only  $\sim 3\%$  memory budget compared to vanilla inference. This is 58%-200% gain compared to SOTA baselines.  
 - Advisor: Dr. Shubham Sharma
- **Snap Inc.** June 2024 - November 2024  
*Research Intern*  
 - Developing backward-compatible training techniques based on hyperbolic geometry to improve performance and efficiency of large retrieval systems, particularly in the context of continual updates of the embedding models.  
 - Advisor: Dr. Tong Zhao
- **VinAI Research (now acquired by Qualcomm)** August 2021 - July 2023  
*Research Resident*  
 - Focusing on robust & trustworthy ML, studying different paradigms of explanation methods for machine learning models and their robustness.  
 - Advisor: Prof. Viet Anh Nguyen  
 - Applied Rotation Project: Interactive Tool for 3D Point Cloud Segmentation.
- **IBM** July 2019 - October 2019  
*AI Research Intern*  
 - Applying PowerAI Vision to visual inspection problems in the car manufacturing process to detect dirt, and dust defects in the car body after painting.  
 - Advisor: Dr. Tam Le Nhan

## AWARDS & HONORS

---



- Honorable Mention in INFORMS Undergraduate Operations Research Prize. 2022
- Best Thesis Presentation Award. 2021
- Problem Winner in ASEAN-India Hackathon. 2021
- Third prize in ACM/ICPC Asia - Ho Chi Minh Regional. 2017
- Third prize in Vietnam Olympiad in Informatics. 2016

## TEACHING ASSISTANT

---

- Trustworthy Deep Learning, *Yale* Spring 2024/Fall 2025
- Applied Algorithms classes, *HUST* 2019 - 2021

## PROFESSIONAL SERVICES

---

- Reviewer at AISTATS 2022/2023, FaCCT 2023, UAI 2023, ICML 2025, NeuRIPS 2023/2024/2025 (Top Reviewer), ICLR 2024/2025.
- Reviewer at ACM Transactions on Sensor Networks (TOSN).
- Organizer at Non-Euclidean Foundation Models and Geometric Learning (NEGEL) Workshop at NeurIPS 2025.